

Project Plan

PINN_M3DC1

Physics-informed surrogate architecture for M3D-C1-shaped nonlinear resistive-MHD workflows

Purpose. This is the living roadmap for PINN_M3DC1. The project begins with a verified two-dimensional reduced-MHD benchmark and a fair data-only-versus-PINN comparison. Genuine M3D-C1 exports enter only through a later adapter and require a new validation claim.

Current status. Version 0.1 establishes the repository backbone, seven-module dependency registry, parallel Python/MATLAB structure, scientific claim hierarchy, initial configuration, and reduced-MHD sign conventions. Textbook-style physics notes are complete for Modules 1, 4, and 5. Modules 4 and 5 freeze the controlled data-only-versus-PINN experiment, including the strong-form residuals, condition and global losses, derivative-order requirements, training curriculum, and fairness contract. No numerical MHD solver or trained surrogate is yet claimed.

Contents

1	Objective and claim boundary	2
2	Initial physics benchmark	2
3	Module roadmap	2
4	Data and model contract	3
5	Controlled comparison	3
6	Verification and validation	3
7	Implementation sequence	4
8	Acceptance milestones	4

1 Objective and claim boundary

The objective is to determine whether physics-informed learning improves a parameterized neural surrogate for nonlinear resistive-MHD evolution under controlled data budgets. Version 1 uses independently generated synthetic arrays. It supports the claim that a PINN can approximate a stated reduced model; it does not establish fidelity to M3D-C1. A validated M3D-C1 surrogate requires genuine code outputs, complete metadata, held-out M3D-C1 cases, and renewed verification.

2 Initial physics benchmark

The periodic two-field model evolves magnetic flux ψ and stream function U . Define

$$[f, g] = f_x g_y - f_y g_x, \quad J = -\nabla^2 \psi, \quad \omega = \nabla^2 U. \quad (2.1)$$

The frozen version-1 residuals are

$$R_\psi = \partial_t \psi + [U, \psi] - \eta \nabla^2 \psi, \quad (2.2)$$

$$R_\omega = \partial_t \omega + [U, \omega] - [J, \psi] - \nu \nabla^2 \omega. \quad (2.3)$$

The first nonlinear case is magnetic-island coalescence. It is small enough for independent implementation while retaining advection, current-sheet formation, resistive reconnection, dissipation, and parameter-dependent time evolution.

3 Module roadmap

No.	Module	Purpose
1	Reduced Resistive-MHD Physics and Conventions	Freeze variables, signs, normalization, boundary conditions, invariants, and diagnostic definitions.
2	Synthetic MHD Solver and Case Generation	Generate verified periodic island-coalescence trajectories and convergence evidence.
3	Dataset Contracts and M3D-C1-Shaped I/O	Store fields, parameters, coordinates, provenance, splits, and normalization in versioned schemas.
4	Data-Only Surrogate Baseline	Establish the fair parameterized neural baseline on complete held-out parameter cases.
5	Physics-Informed Surrogate	Add PDE, initial, boundary, and optional global-balance constraints to the same learning task.
6	Validation, Diagnostics, and Benchmarking	Compare field accuracy, physical histories, residuals, generalization, runtime, and ablations.
7	M3D-C1 Export Adapter and Transfer Learning	Map genuine exported fields and metadata into the project contract without silently changing physical meaning.

4 Data and model contract

A synthetic case stores (x, y, t) , ψ , U , J , ω , η , ν , A_0 , normalization, discretization, boundary conditions, solver provenance, and diagnostic histories. The initial surrogate is

$$(x, y, t, \eta, \nu, A_0) \mapsto (\hat{\psi}, \hat{U}), \quad (4.1)$$

with $\hat{J}$ and $\hat{\omega}$ derived by automatic differentiation. Plot images are not training data because they omit the numerical field values and metadata needed to define this map.

5 Controlled comparison

Model	Data loss	Physics losses	Role
Reduced-MHD solver	–	Discretized equations	Synthetic reference
Data-only network	yes	no	Fair neural baseline
Parameterized PINN	yes	PDE/IC/BC	Physics-information test

The two neural models use identical case splits, observations, and comparable capacity. Entire (η, ν, A_0) combinations are excluded from training. Interpolation and extrapolation are reported separately.

6 Verification and validation

The reference solver must pass derivative, bracket, periodicity, divergence, convergence, and energy-behavior checks before producing training data. Surrogate evaluation includes relative L_2 and L_∞ errors for all fields, peak-current and current-sheet errors, magnetic and kinetic energy histories, reconnection metrics, PDE residuals, and periodicity. Runtime reporting includes reference generation, training, inference, and break-even query count.

7 Implementation sequence

1. Freeze equations, normalization, case schema, and diagnostic definitions.
2. Implement the independent pseudo-spectral solver in Python and MATLAB.
3. Produce one converged island-coalescence reference trajectory.
4. Generate a parameter sweep and freeze whole-case data splits.
5. Train the data-only parameterized baseline.
6. Debug a single-case PINN, then train the parameterized PINN.
7. Run ablations, generalization tests, uncertainty checks, and cost analysis.
8. Add and validate an adapter only when genuine M3D-C1 arrays are available.

8 Acceptance milestones

M1 – synthetic solver. One command writes a reproducible HDF5 case and diagnostics; refinement and dissipation checks pass.

M2 – fair learning experiment. Data-only and PINN models share cases, samples, capacity assumptions, and scoring.

M3 – held-out validation. Field and physics metrics are reported on unseen parameter combinations, with interpolation separated from extrapolation.

M4 – M3D-C1 transfer. Exported variables, mesh, normalization, closures, and boundary conditions are mapped explicitly; the full test suite is rerun before an M3D-C1 surrogate claim is made.